Distributed Transactions

Dominic Duggan

Stevens Institute of Technology

Based in part on materials by K. Birman

1

Client, server, and databases

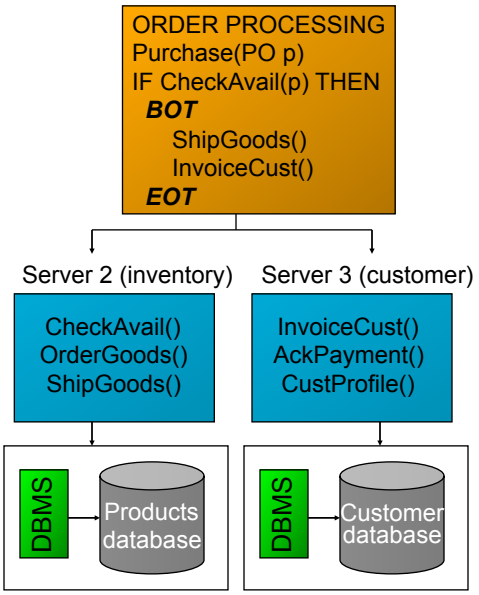

2

Multiple Databases

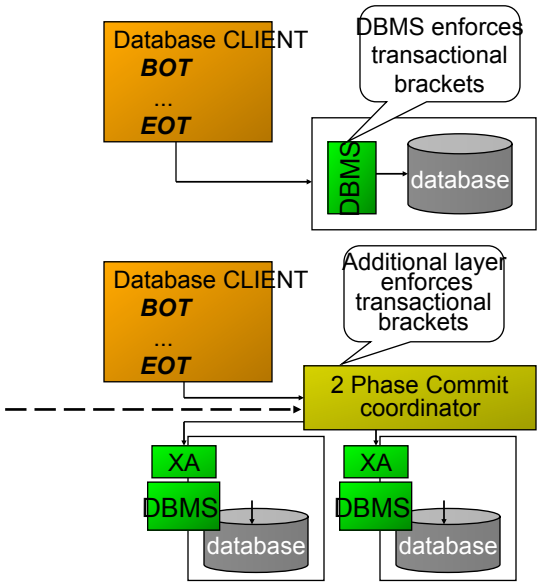

3

Transactional RPC

- Transaction Managers
- Transaction Processing Monitor (TPM)
- XA API

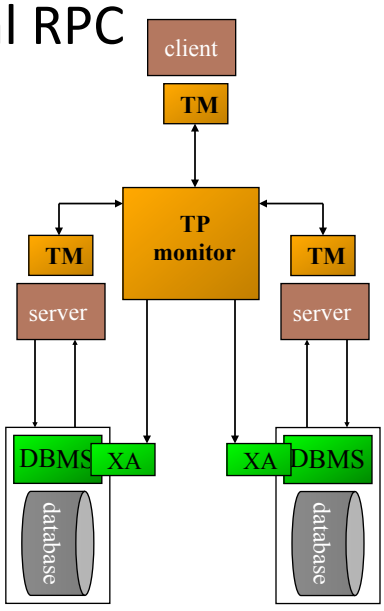

4

Basic TRPC (making calls)

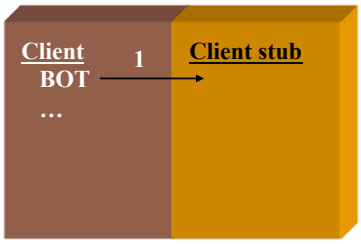

5

Basic TRPC (making calls)

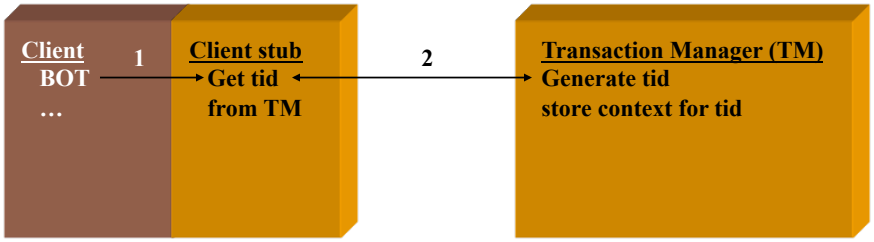

6

Basic TRPC (making calls)

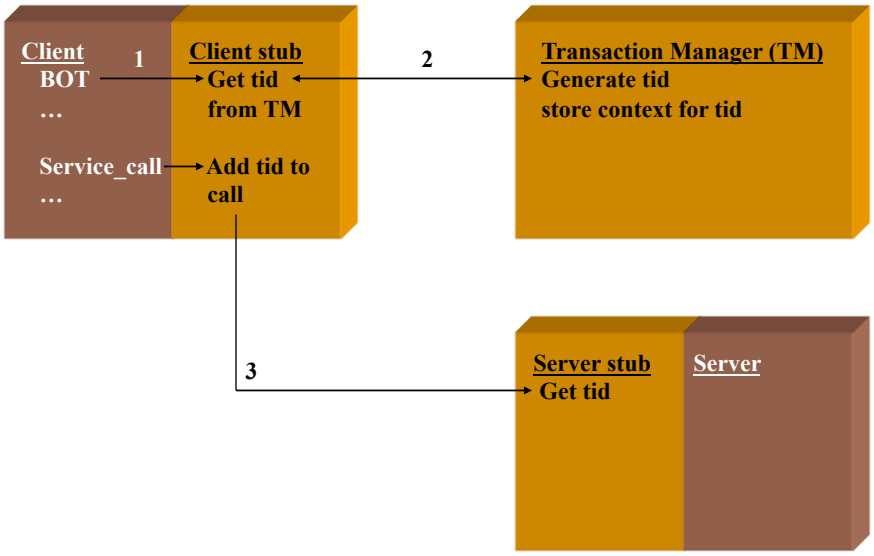

7

Basic TRPC (making calls)

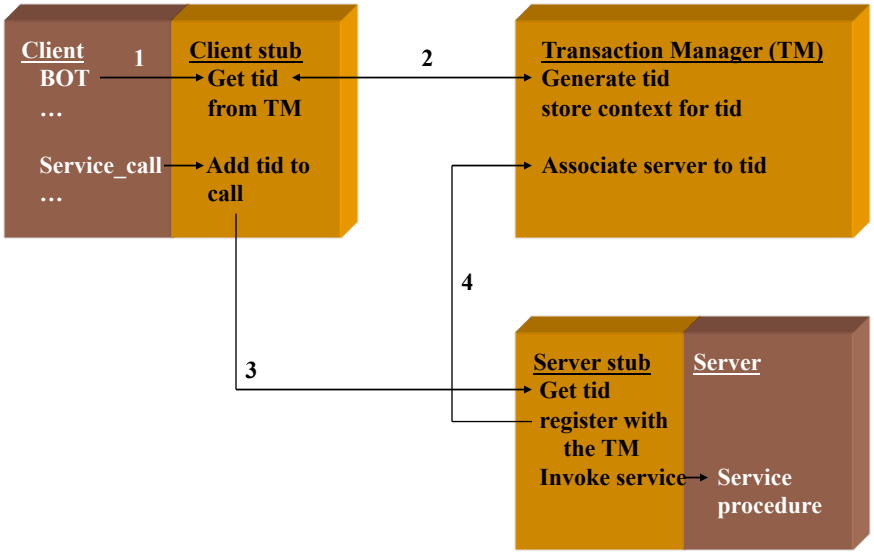

8

Basic TRPC (making calls)

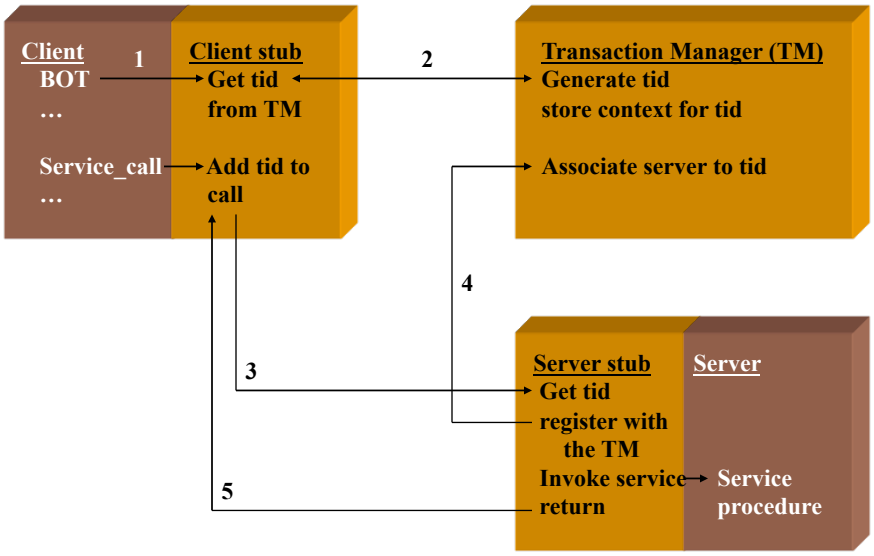

9

Basic TRPC (committing calls)

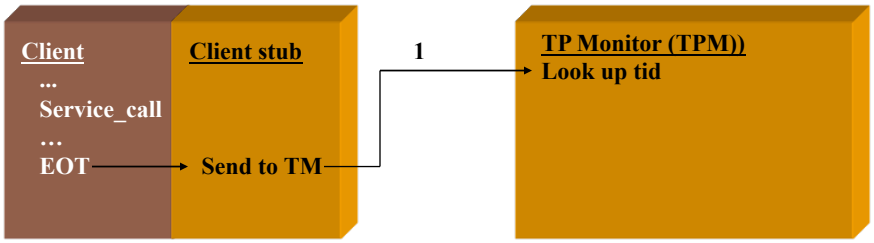

10

Basic TRPC (committing calls)

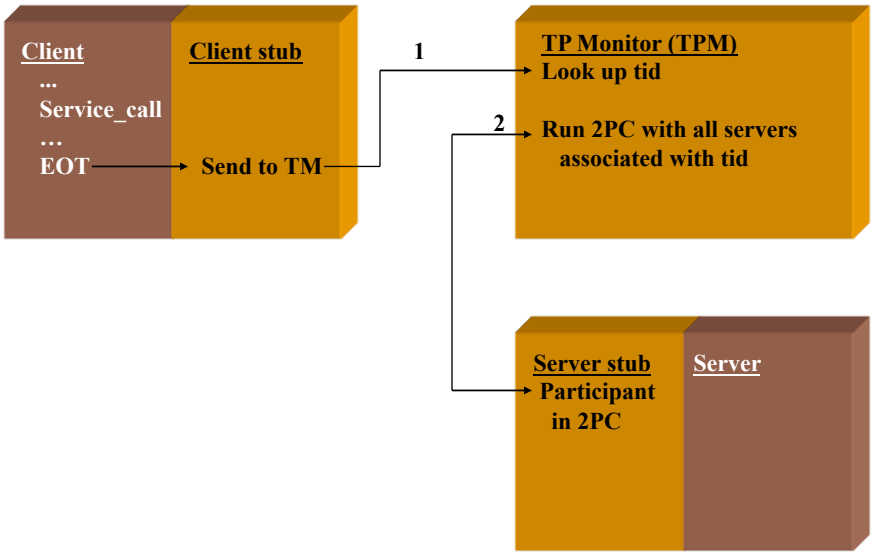

11

Basic TRPC (committing calls)

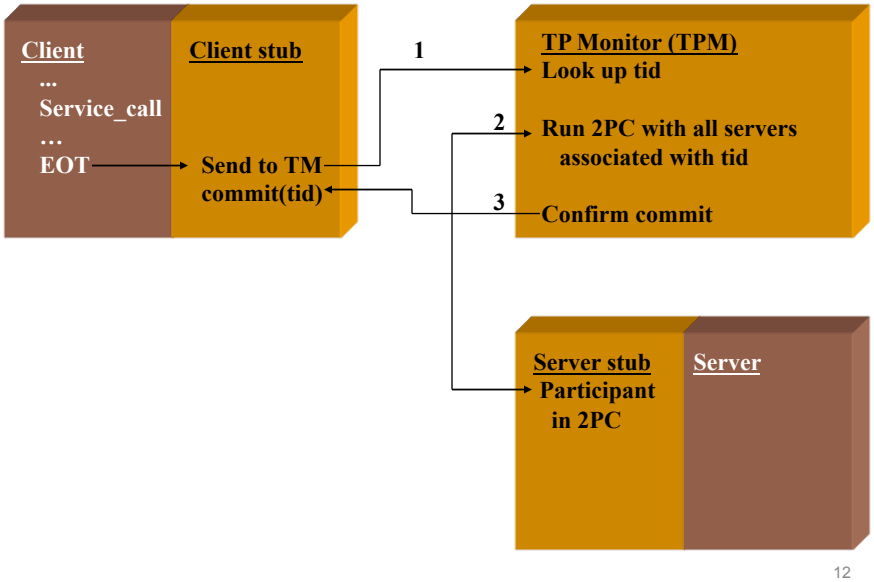

12

TWO PHASE COMMIT (2PC)

13

Atomic Commitment

- Given a set of processes
- **Coordinator** (i.e. TPM) wants to initiate an action (commit)
- **Participants** may vote for or against the action
- Perform the action only if all vote in favor
- Otherwise abort
- Goal is *all-or-nothing* outcome

14

Non-triviality

- Avoid solutions that do nothing
- What is a trivial solution?
- What is a validity condition?

15

Commit protocol illustrated

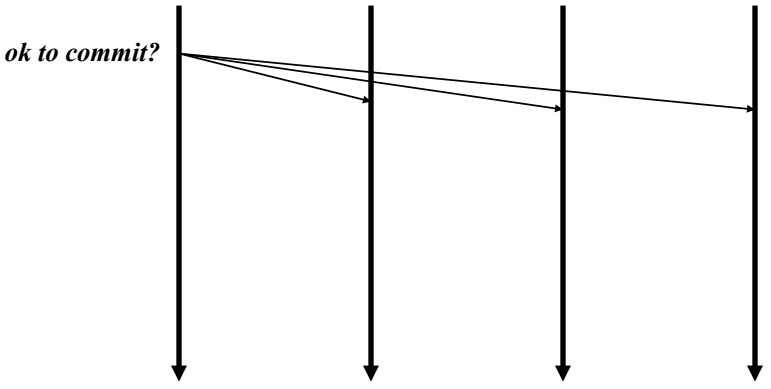

16

Commit protocol illustrated

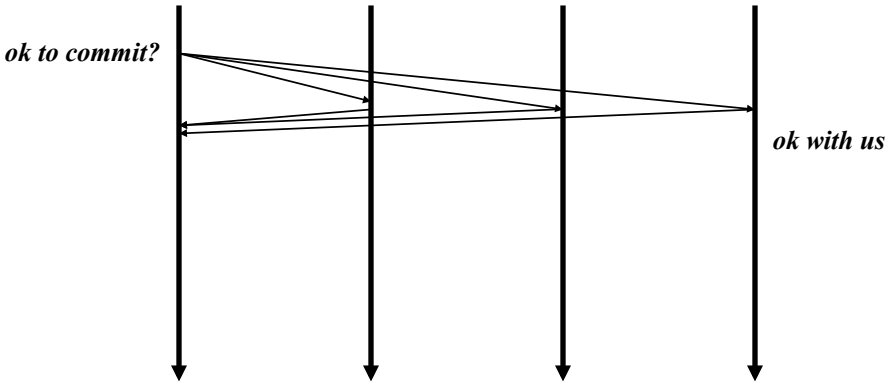

17

Commit protocol illustrated

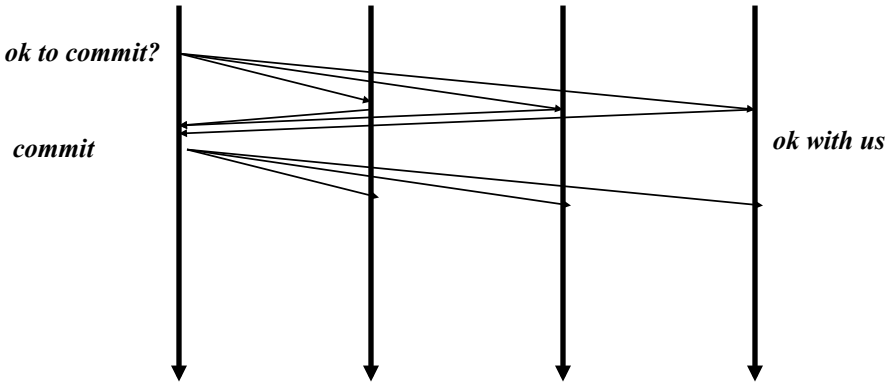

Note: garbage collection protocol not shown here

18

Two-phase Commit

- Phase 0: Flush caches on Web, app server
- Phase 1: Coordinator asks participants for vote
  - Data managers force updates to the log
  - Then say "ok to commit"
- Phase 2: If all are ok to commit, then coordinator tells participants to commit. Otherwise, abort.
- Data managers then make updates permanent or rollback to old values, and release locks

19

Commit protocol illustrated

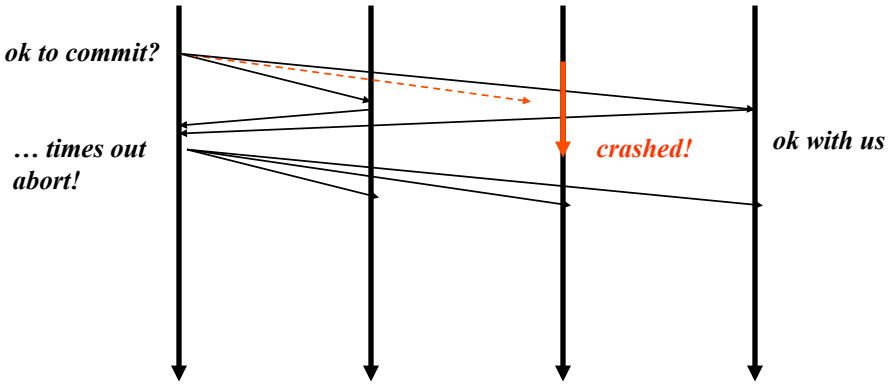

Note: garbage collection protocol not shown here

20

Non-triviality

- Avoid solutions that do nothing
- Commit validity: if all vote for commit, protocol must commit
- …but what if participant vote is lost?
- "Non-triviality" condition hard to capture

21

Unilateral abort

- Any data manager can unilaterally abort a transaction until it has said "prepared"
- Implication: even a data manager where only reads were done must participate in 2PC protocol!

22

PROBLEMS WITH 2PC

Non-blocking Commit

- Goal: a protocol that allows all operational processes to terminate the protocol even if some subset crash

Commit with unreliable failure detectors

- Assume processes fail by crashing
  - No Byzantine failures
- Coordinator detects failures (unreliably) using timouts
- Challenge: terminate the protocol if the coordinator fails

25

2 phase commit

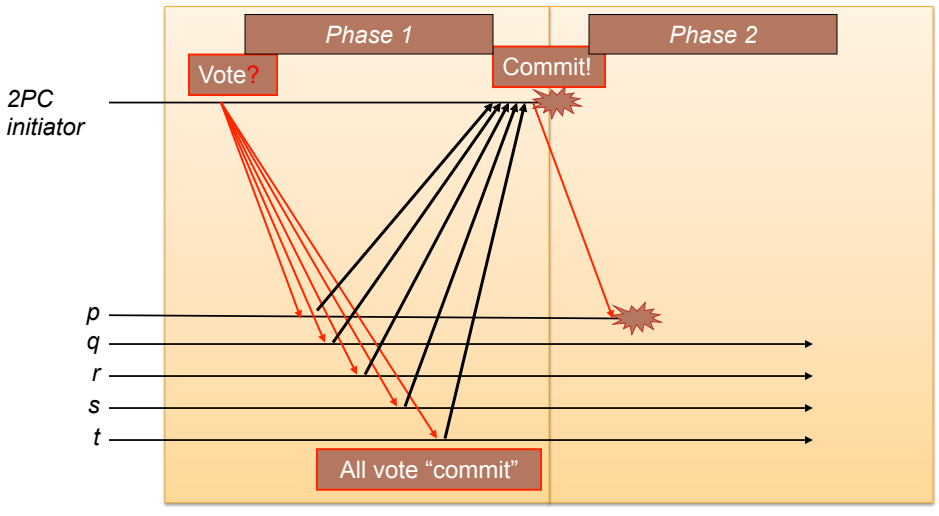

THREE PHASE COMMIT (3PC)

27

Three-phase commit

- Seeks to increase availability
- Makes an unrealistic assumption that failures are accurately detectable
- With this, can terminate the protocol even if a failure does occur

28

Three-phase commit

- Coordinator starts protocol by sending request
- Participants vote to commit or to abort
- Coordinator collects votes, decides on outcome
- Coordinator can abort immediately
- To commit, coordinator first sends a "prepare to commit" message
- Participants acknowledge, commit occurs during a final round of "commit" messages

29

Three-phase commit

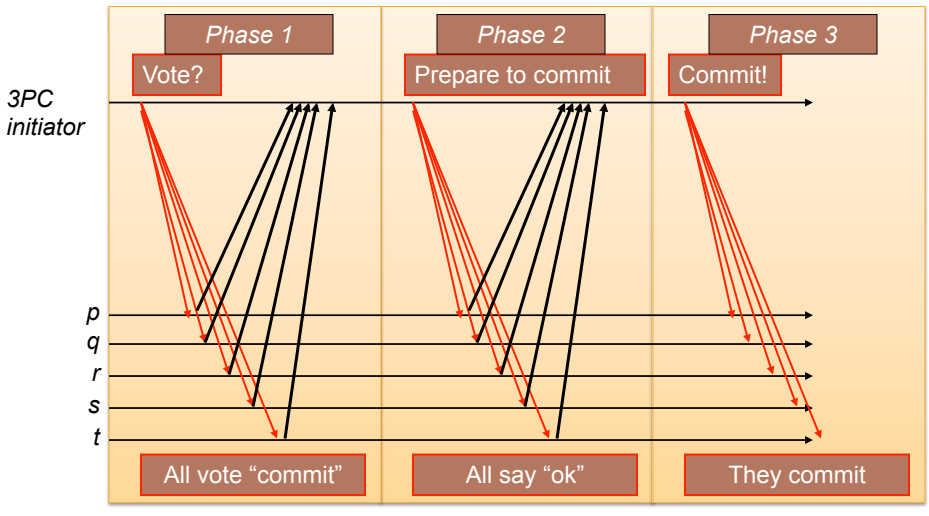

State diagram for non-faulty member

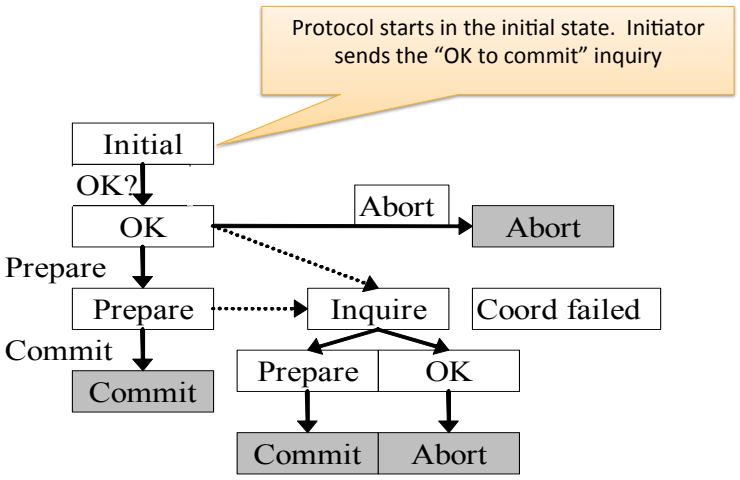

31

State diagram for non-faulty member

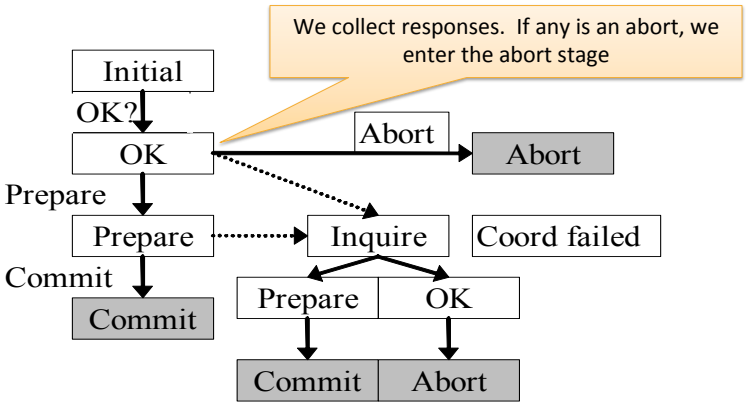

32

State diagram for non-faulty member

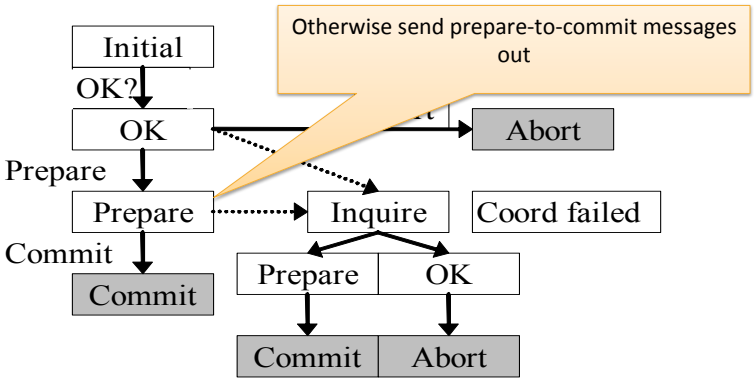

33

State diagram for non-faulty member

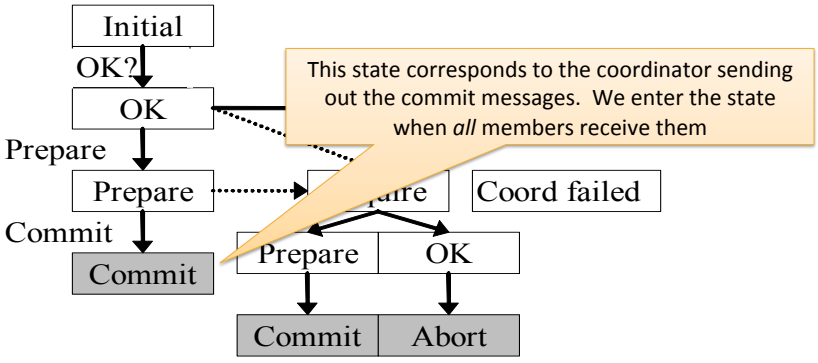

34

State diagram for non-faulty member

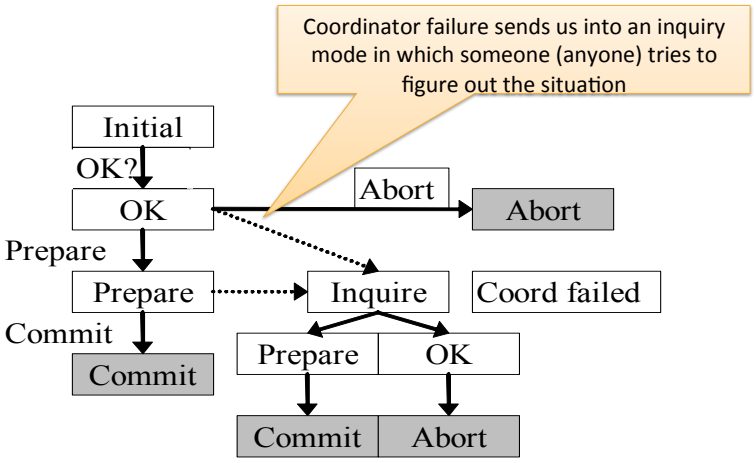

35

State diagram for non-faulty member

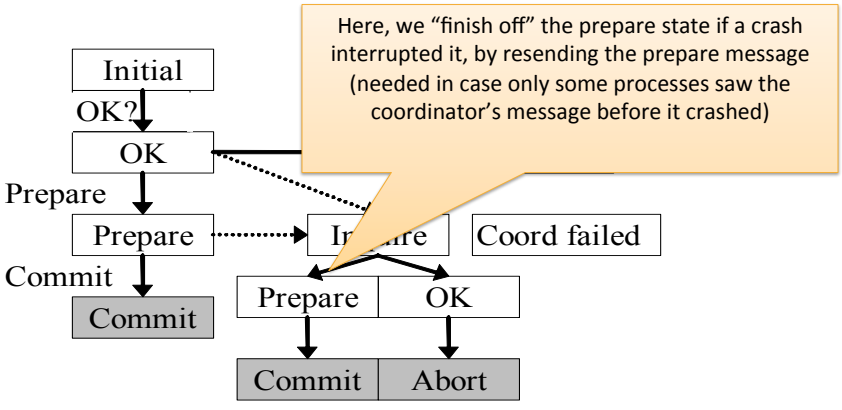

36

State diagram for non-faulty member

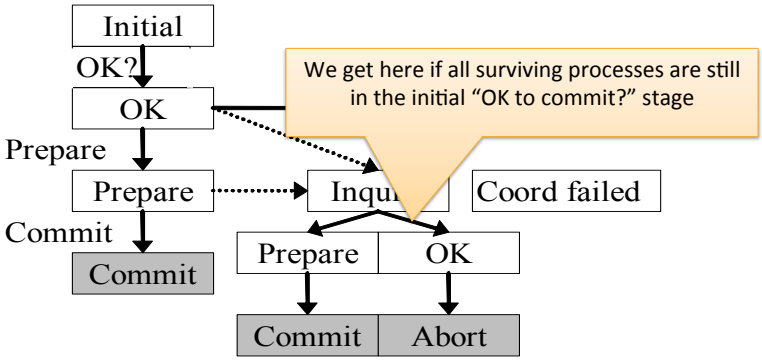

37

State diagram for non-faulty member

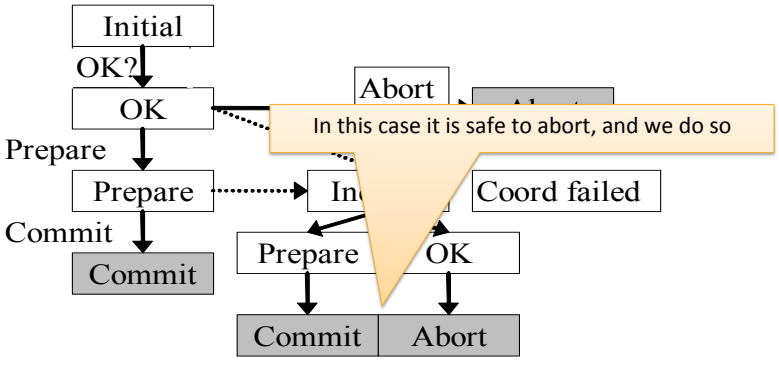

38

Observations about 3PC

- Key point: Extra buffer state
- What if none of surviving participants have heard from coordinator?
  - After voting phase
  - 2PC: Some crashed processes may have committed
  - 3PC: No crashed process has committed yet (Why?)

39

Observations about 3PC

- If any process is in "prepare to commit" all voted for commit
- Protocol commits only when all surviving processes have acknowledged prepare to commit
- After coordinator fails, it is easy to run the protocol forward to commit state (or back to abort state)

40

Problems with 3PC

- Assumes reliable failure detectors
- But even with realistic failure detectors (that can make mistakes), protocol still blocks!
  - "Network partitioning"
- Can prove that this problem is not avoidable

41

Situation in practical systems?

- Most use protocols based on 2PC
- Need to extend garbage collection
  - protocol state information
- Some systems accept the risk of blocking
- Others reduce the consistency property to make progress

42

Example: Primary with Backup

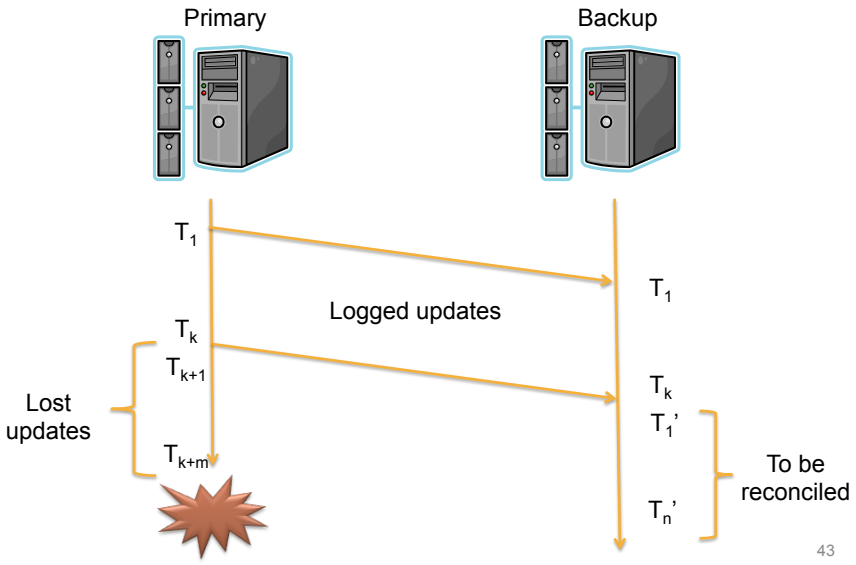

43

PAXOS AND ATOMIC COMMIT

44

Database Commit

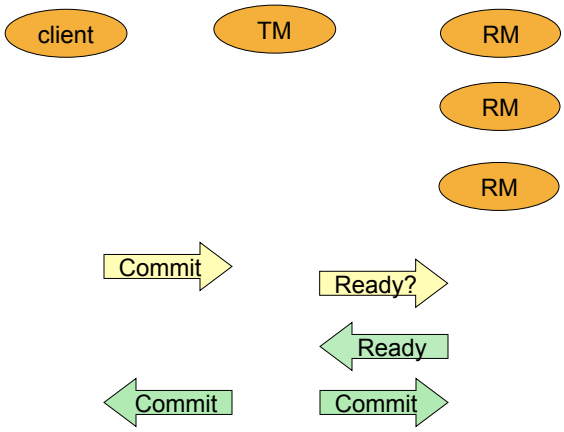

TM: Transaction Manager
RM: Resource Manager

45

Two Phase Commit

- *N* Resource Managers (RMs)
- Want all RMs to commit or all abort.
- Coordinated by Transaction Manager (TM)
  TM sends Prepare, Commit-Abort
- RM responds Prepared, Aborted
- *3N+1* messages
- *N+1* stable writes
- Delay
  - 4 message
  - 2 stable write
- Blocking:
  if TM fails,
  Commit-Abort stalls

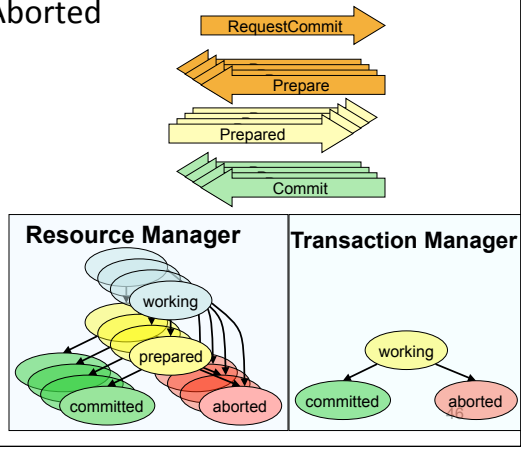

The Problem With 2PC

- Atomicity – all or nothing
- Consistency – does right thing
- Isolation – no concurrency anomalies
- Durability / Reliability – state survives failures
- Availability: always up

Blocks if TM fails

47

Problem Statement

- ACID Transactions make error handling easy.
- One fault can make 2-Phase Commit block.
- Goal: ACID and Available.
  Non-blocking despite *F* faults.

48

Fault-Tolerant Two Phase Commit

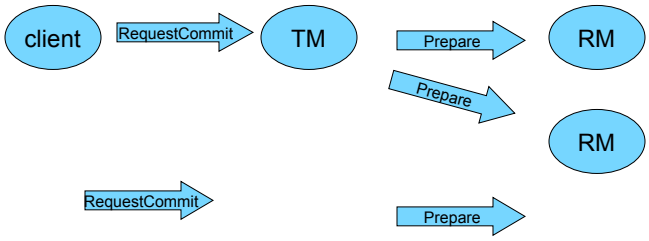

49

Fault-Tolerant Two Phase Commit

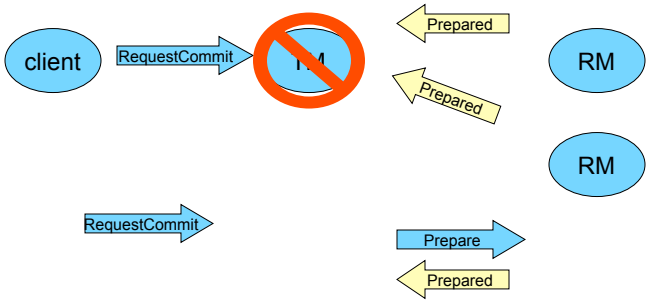

If the 2PC Transaction Manager (TM) Fails, transaction blocks.

50

Fault-Tolerant Two Phase Commit

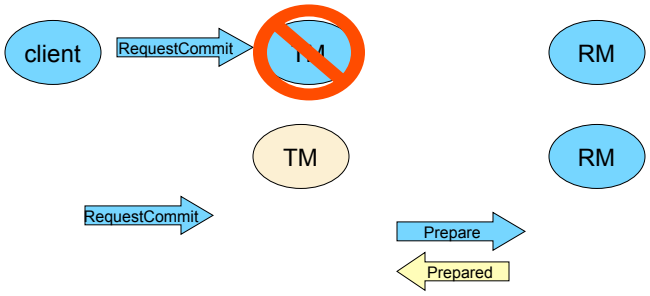

If the 2PC Transaction Manager (TM) Fails, transaction blocks.

Solution: Add a “spare” transaction manager
(non blocking commit, 3 phase commit)

51

Fault-Tolerant Two Phase Commit

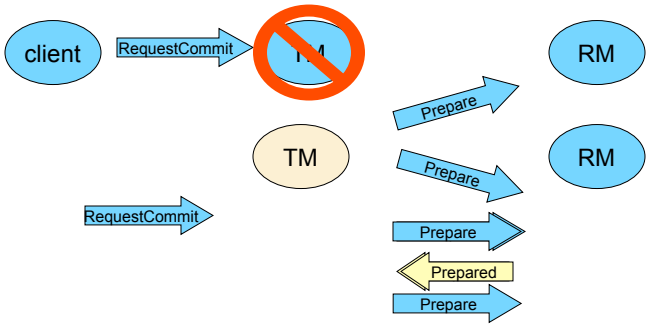

If the 2PC Transaction Manager (TM) Fails, transaction blocks.

Solution: Add a “spare” transaction manager
(non blocking commit, 3 phase commit)

52

Fault-Tolerant Two Phase Commit

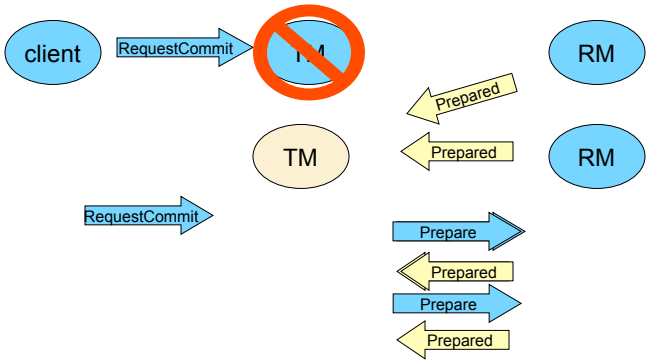

If the 2PC Transaction Manager (TM) Fails, transaction blocks.

Solution: Add a “spare” transaction manager
(non blocking commit, 3 phase commit)

53

Fault-Tolerant Two Phase Commit

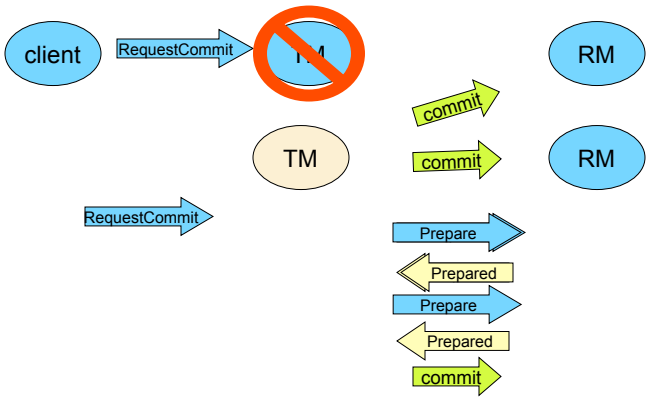

If the 2PC Transaction Manager (TM) Fails, transaction blocks.

Solution: Add a “spare” transaction manager
(non blocking commit, 3 phase commit)

54

Fault-Tolerant Two Phase Commit

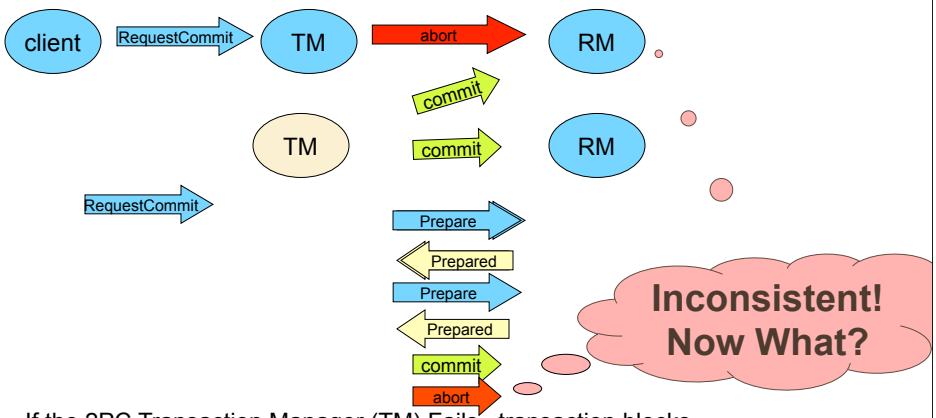

If the 2PC Transaction Manager (TM) Fails, transaction blocks.

Solution: Add a "spare" transaction manager
(non blocking commit, 3 phase commit)

55

Paxos Consensus Box

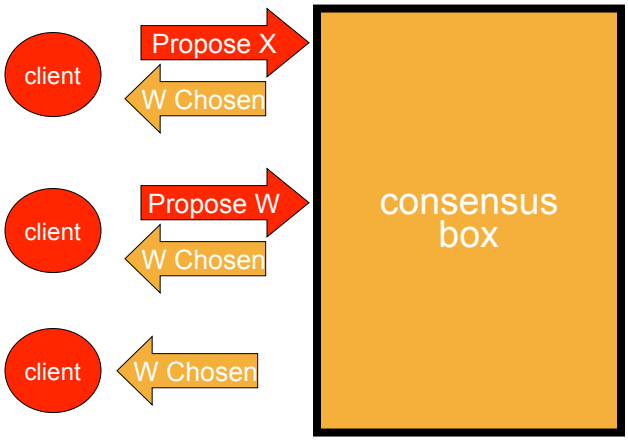

- Collects proposed values
- Picks one proposed value
- Remembers it forever

56

Consensus for Commit
The Obvious Approach

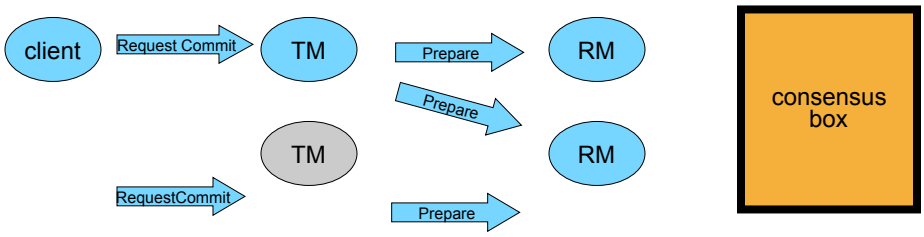

- Get consensus on TM's decision.
- TM just learns consensus value.
- TM is "stateless"

57

Consensus for Commit
The Obvious Approach

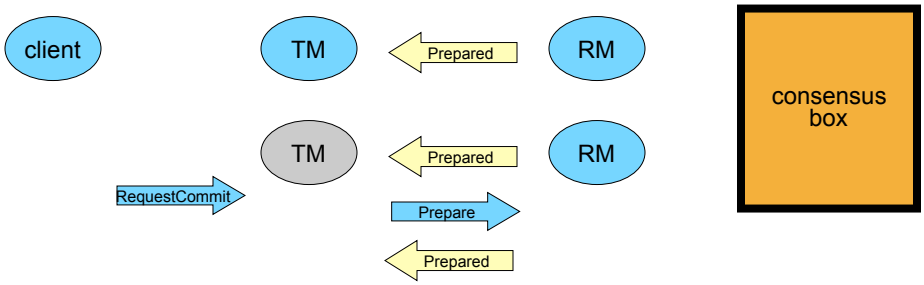

- Get consensus on TM's decision.
- TM just learns consensus value.
- TM is "stateless"

58

Consensus for Commit
The Obvious Approach

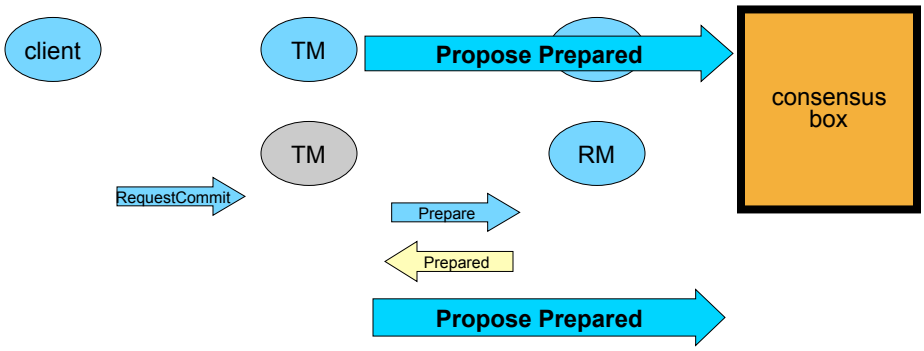

- Get consensus on TM's decision.
- TM just learns consensus value.
- TM is "stateless"

59

Consensus for Commit
The Obvious Approach

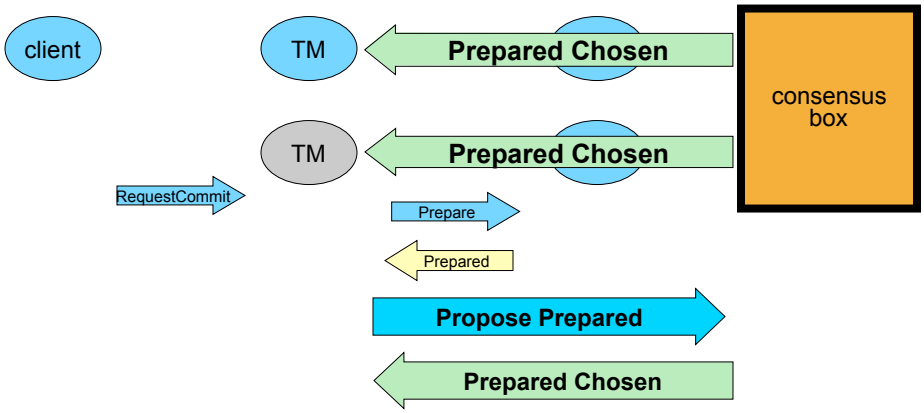

- Get consensus on TM's decision.
- TM just learns consensus value.
- TM is "stateless"

60

Consensus for Commit
The Obvious Approach

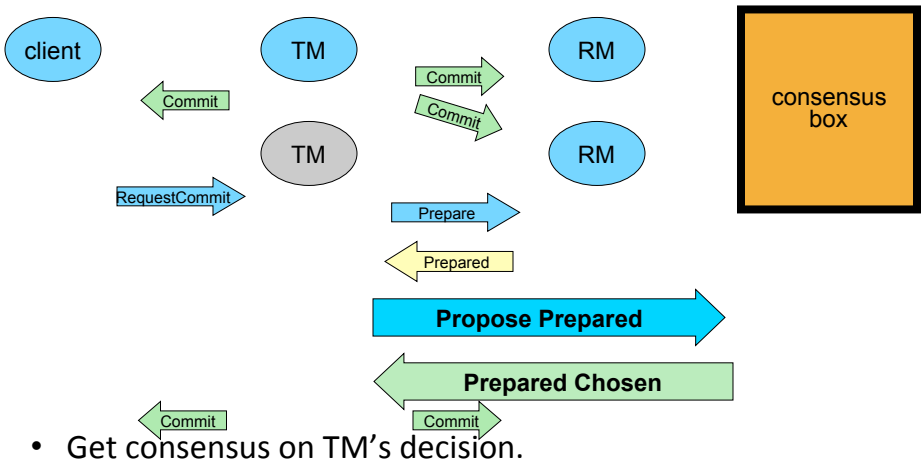

- Get consensus on TM's decision.
- TM just learns consensus value.
- TM is "stateless"

61

Consensus for Commit
The Paxos Commit Approach

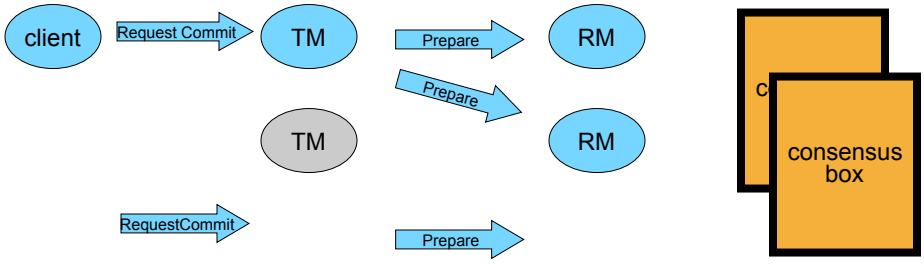

- Get consensus on each RM's choice.
- TM just combines consensus values.
- TM is "stateless"

62

Consensus for Commit
The Paxos Commit Approach

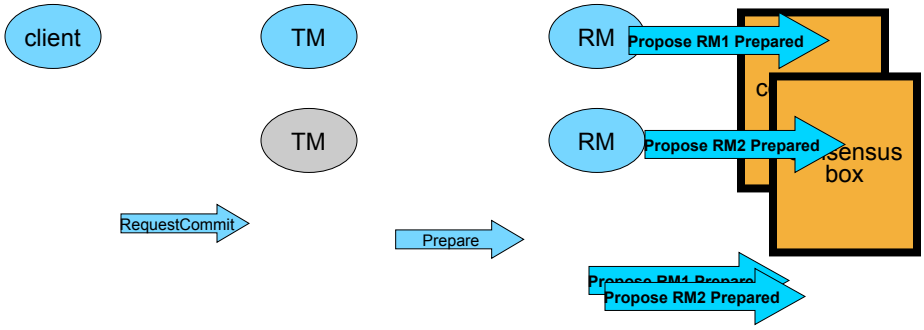

- Get consensus on each RM's choice.
- TM just combines consensus values.
- TM is "stateless"

63

Consensus for Commit
The Paxos Commit Approach

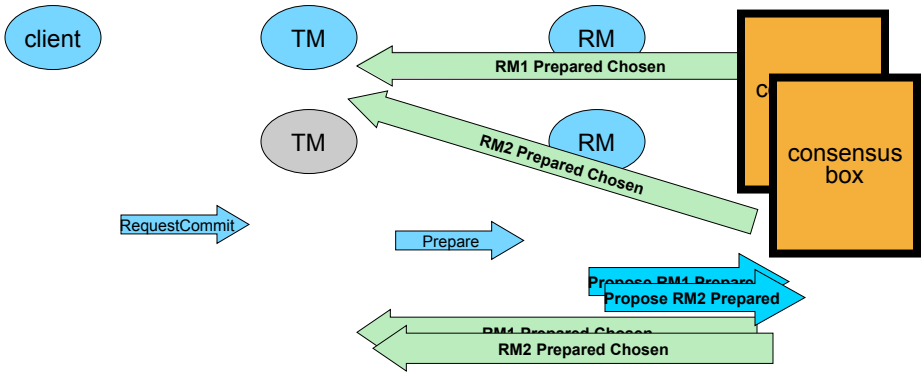

- Get consensus on each RM's choice.
- TM just combines consensus values.
- TM is "stateless"

64

Consensus for Commit
The Paxos Commit Approach

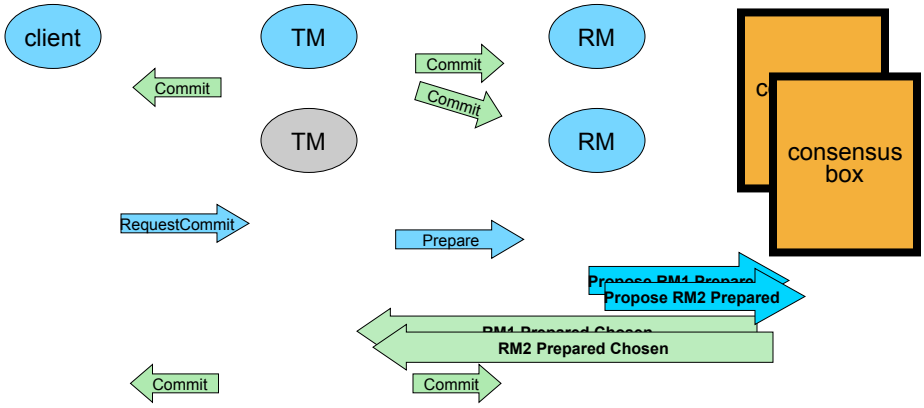

- Get consensus on each RM's choice.
- TM just combines consensus values.
- TM is "stateless"

Consensus in Action

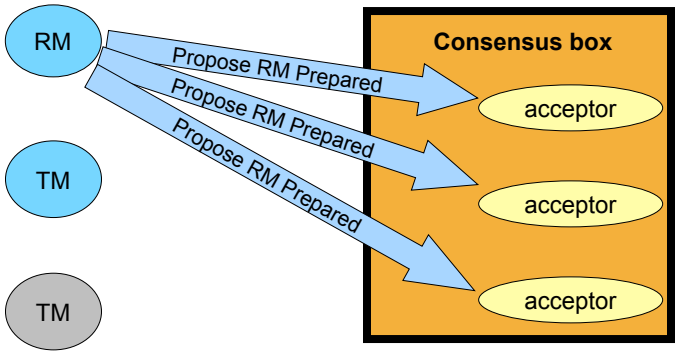

- The normal (failure-free) case

Consensus in Action

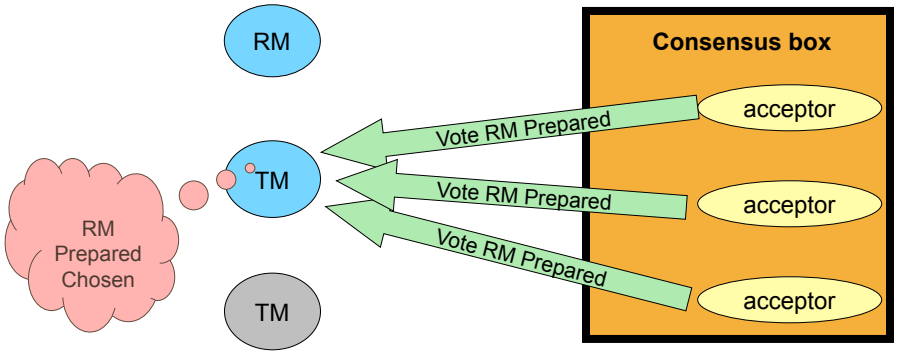

- The normal (failure-free) case
- Two message delays
- Can optimize

67

Consensus in Action

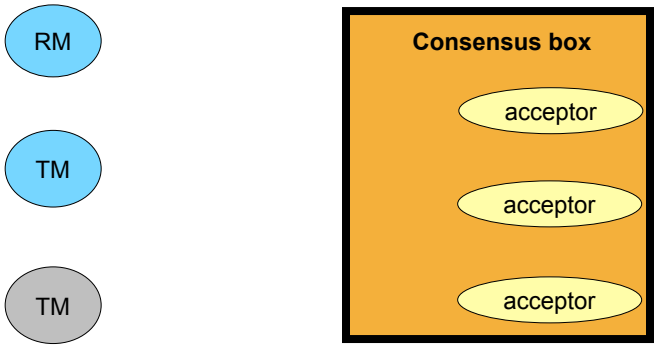

TM can always learn what was chosen,
or get *Aborted* chosen if nothing chosen yet;

68

Consensus in Action

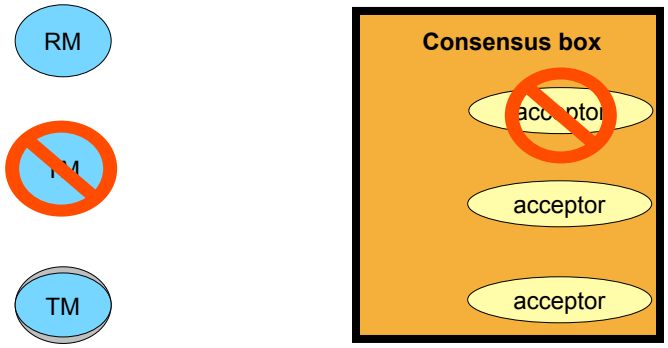

TM can always learn what was chosen,
or get *Aborted* chosen if nothing chosen yet;
if majority of acceptors working .

69

The Obvious Approach

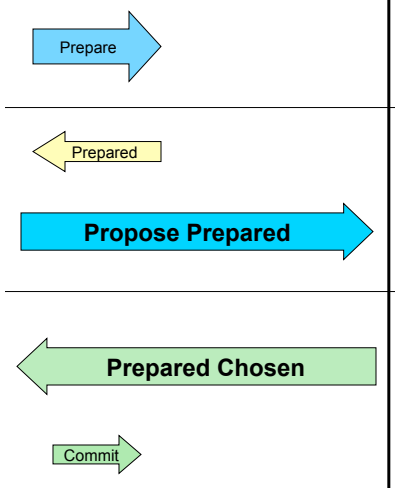

Paxos Commit

One fewer message delay

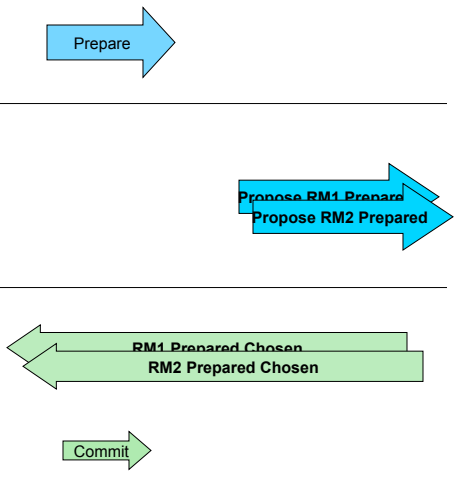

70

Paxos Commit

- *N* RMs
- *2F+1* acceptors (~*2F+1* TMs)
- If *F+1* acceptors see all RMs prepared, then transaction committed.
- *(N+1)(F+3) - 2* messages
  - 1 request to commit
  - N - 1 prepare
  - N (F + 1) prepared
  - F + 1 all prepared
  - N commit
- 5 message delays
  2 stable write delays.

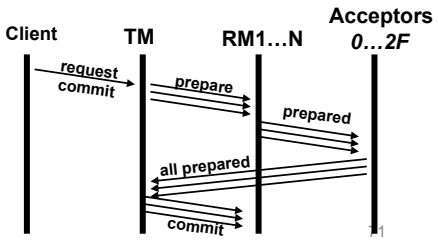

71

Two-Phase Commit	Paxos Commit
	tolerates *F* faults
• 3*N*+1 messages	• 3*N*+ 2*F*(*N*+1) +1 messages
• *N*+1 stable writes	• *N*+2*F*+1 stable writes
• 4 message delays	• 5 message delays
• 2 stable-write delays	• 2 stable-write delays

Same algorithm when *F*=0 and TM = Acceptor

72

EXTENDED TRANSACTION MODELS

73

Savepoints

- Idea: allow the state of a transaction to be rolled back to a certain point in execution
- Not necessary to roll back the entire transaction

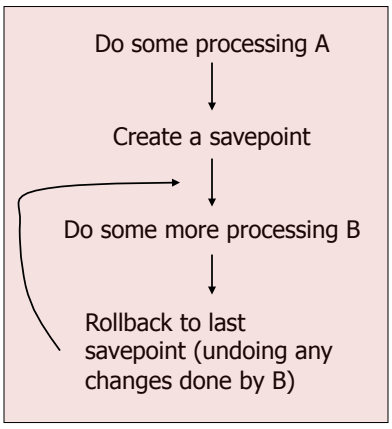

74

Nested Transactions (1)

- Idea: Allow transactions to be nested inside other transactions
- Child transaction's changes not visible to parent, siblings until commit
- Failure of a child transaction does not imply failure of the parent

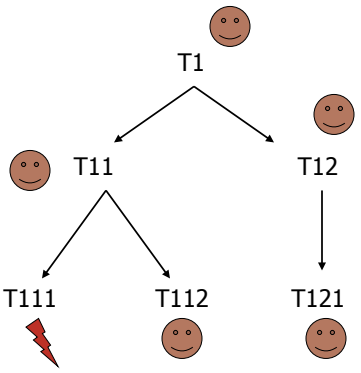

75

Nested Transactions (2)

- Abort of a transaction forces all of its descendants to abort
- So commit of a non-root transaction is *tentative*
- Child can obtain locks from parent (*lock inheritance*)
- Child's locks released to parent (*lock anti-inheritance*)

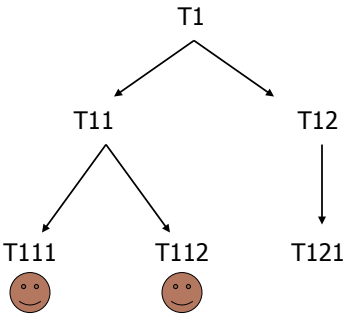

76

Nested Transactions (2)

- Abort of a transaction forces all of its descendants to abort
- So commit of a non-root transaction is *tentative*
- Child can obtain locks from parent (*lock inheritance*)
- Child's locks released to parent (*lock anti-inheritance*)

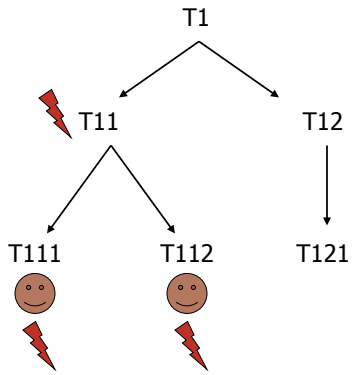

77

Nested Transactions (2)

- Abort of a transaction forces all of its descendants to abort
- So commit of a non-root transaction is *tentative*
- Child can obtain locks from parent (*lock inheritance*)
- Child's locks released to parent (*lock anti-inheritance*)

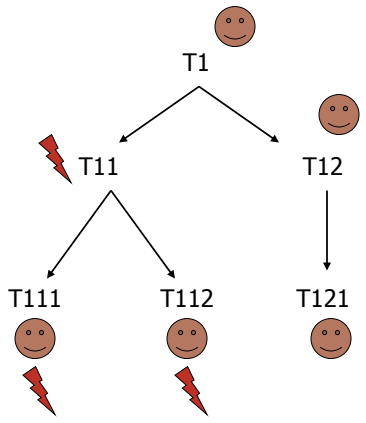

78

Multi-Level Transactions (1)

- Unlike Nested, Multi-Level allows a transaction's changes to be made visible to *everyone* when it commits
  - Nested: visible to parent
- A compensating transaction is run if the parent transaction subsequently aborts
  - Each transaction has its own specific compensating transaction

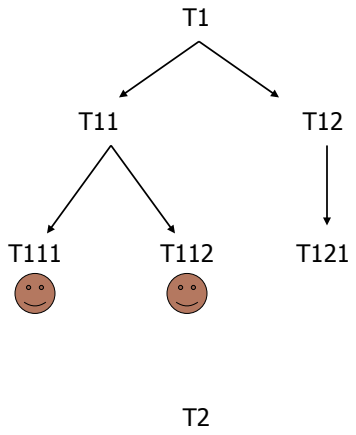

79

Multi-Level Transactions (1)

- Unlike Nested, Multi-Level allows a transaction's changes to be made visible to *everyone* when it commits
  - Nested: visible to parent
- A compensating transaction is run if the parent transaction subsequently aborts
  - Each transaction has its own specific compensating transaction

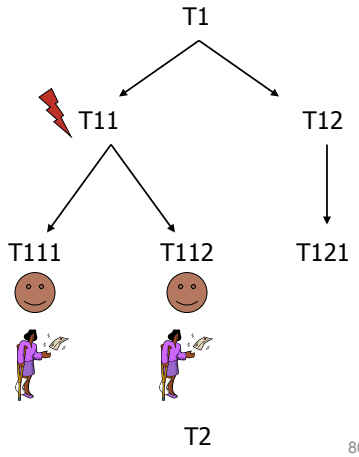

80

Multi-Level Transactions (1)

- Unlike Nested, Multi-Level allows a transaction's changes to be made visible to *everyone* when it commits
  - Nested: visible to parent
- A compensating transaction is run if the parent transaction subsequently aborts
  - Each transaction has its own specific compensating transaction

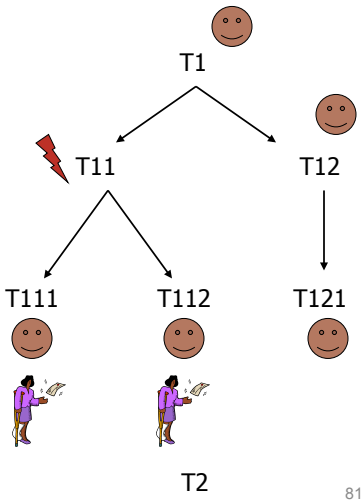

81

Multi-Level Transactions (2)

- Motivation for multi-level: sharing of low-level resources in layered software systems
- Ex: insert records into database:
  - Find page on disk with free space
  - Insert records
  - Update indexes
  - Commit
- Nested would require page to be locked to further allocation until commit

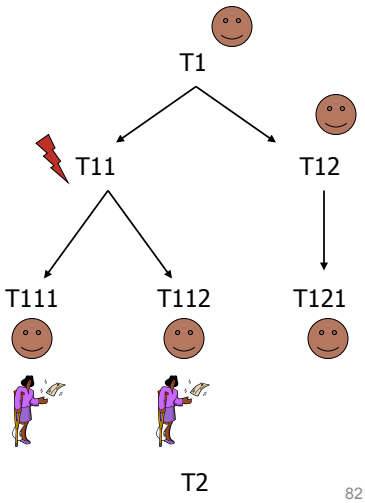

82

Sagas (1)

- Transactions are incompatible with long-lived applications
  - Locks must be held to ensure ACID properties
- Sagas relax ACID properties, allow intermediate states to be visible to outside
- A saga is a sequence of transactions
  - "long-lived transaction"

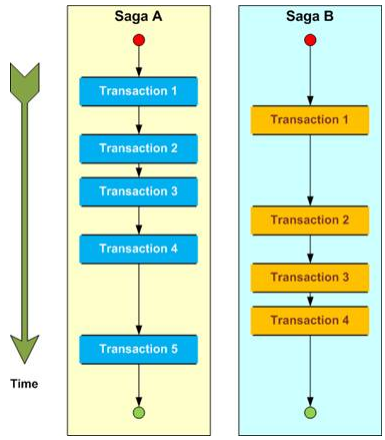

83

Sagas (2)

- Each transaction T_i has a corresponding compensating transaction CT_i
- If T_k aborts in a transaction, then the compensations

  $CT_{k-1},\ldots,CT_2,CT_1$

  are executed, in that order
  - Explicit rollback

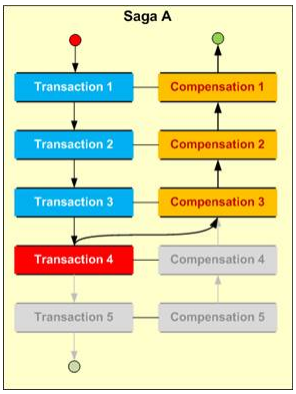

84

Sagas (3)

- Sagas also support forward recovery using savepoints
- If state is persisted between transactions, then compensating transactions roll back to that savepoint
- Saga continues from that point forward

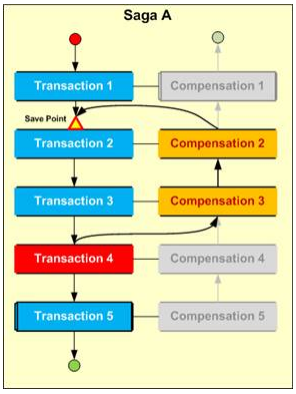

85